

GULAM MOHAMMED MUBARIZUDDIN

Data Engineer | AWS | Snowflake | Airflow | Python | SQL

Hyderabad, Telangana, India | gulammohammedmubarizuddin@gmail.com | [LinkedIn](#) | [GitHub](#) | [Portfolio](#)

PROFESSIONAL SUMMARY

Data Engineer with around 4 years of hands-on experience designing, building, and supporting production cloud data pipelines, specializing in AWS + Snowflake + Apache Airflow architectures. Proven track record delivering end-to-end ingestion, transformation, incremental loading, and orchestration solutions for large-scale content, engagement, and telemetry datasets (100s of GB to multi-TB scale). AWS Certified Solutions Architect – Associate with strong expertise in Snowflake data modeling, SQL & Python transformations, Airflow DAG development, data quality frameworks, and production support. Additional experience with Azure Databricks and PySpark for high-volume data processing.

TECHNICAL SKILLS

Cloud Platforms: AWS, GCP, Azure

AWS Services: Amazon S3, AWS Glue, AWS Lambda, IAM, CloudWatch

Data Warehousing: Snowflake (Tables, Schemas, Incremental Loads, Clustering), Dimensional Modeling, ELT

Orchestration & Processing: Apache Airflow, PySpark, Spark SQL, Databricks, Batch Processing

Programming & Databases: Python, SQL, Linux, Relational Databases

Data Lake / Storage: Amazon S3, Delta Lake, Partitioned Storage, Lakehouse Architecture

Engineering Practices: Data Validation & Quality Checks, Schema Management, Error Handling, Logging, Monitoring, Incremental Processing, Git, Production Support

PROFESSIONAL EXPERIENCE

Wipro | Data Engineer | Hyderabad, India

June 2024 – Present

- Assigned to a US autonomous-mobility client engagement; develop and support cloud-based data pipelines processing multi-TB scale telemetry and platform datasets consumed by analytics and engineering teams.
- Build PySpark and Spark SQL transformation pipelines on Azure Databricks for data ingestion, cleansing, transformation, and preparation of analytical datasets.
- Implement incremental processing and partitioning strategies to efficiently handle large-volume datasets and minimize full reprocessing.
- Develop and schedule 10+ Databricks workflows and notebooks for recurring production workloads with robust validation, logging, error handling, and failure-recovery mechanisms.
- Support schema evolution and production incidents while maintaining data quality and pipeline availability; collaborate with technical and business stakeholders to resolve data issues.
- Recognized with the Wipro Star Performer Award for on-time delivery and quality execution of data platform work.

P.R. Associates (Inshorts) | Data Engineer | Hyderabad, India

June 2023 – June 2024

- Designed, developed, and supported production AWS + Snowflake + Apache Airflow data pipelines for a high-traffic digital content platform, processing 200–400 GB of content, engagement, and metadata data daily.
- Built end-to-end data ingestion workflows: landed source data in Amazon S3, applied Python/SQL transformations, and loaded curated datasets into 40+ Snowflake tables supporting analytics and reporting.
- Implemented incremental loading patterns across key pipelines, reducing daily full-table reprocessing by ~65% and significantly improving overall pipeline runtime and cost efficiency.
- Developed and maintained 15+ Apache Airflow DAGs for reliable orchestration, scheduling, dependency management, monitoring, and alerting of data pipelines.
- Performed systematic data validation and quality checks (completeness, uniqueness, consistency) on incoming feeds, helping maintain >99% data reliability for downstream analytics consumers.
- Standardized Snowflake table definitions and applied basic clustering strategies to support analytical workloads and reduce ad-hoc data cleanup efforts.
- Collaborated with product and data owners to gather requirements, define test cases, validate pipeline outputs, and drive timely resolution of data issues.
- Provided production support including troubleshooting, root-cause analysis, and resolution of data ingestion and transformation failures, contributing to high pipeline availability.

KEY PROJECT EXPERIENCE

AWS + Snowflake Data Engineering Pipeline

Technologies: AWS S3, Snowflake, Apache Airflow, Python, SQL, Linux

- Built a complete production-grade cloud data pipeline ingesting, transforming, and loading 200–400 GB/day of content, engagement, and metadata datasets into Snowflake.

- Used Amazon S3 as the landing and staging layer; implemented transformation and validation logic in Python and SQL before loading into 40+ curated Snowflake tables.
- Orchestrated the entire workflow using 15+ Apache Airflow DAGs with proper dependency management, scheduling, monitoring, and alerting.
- Implemented incremental processing that reduced unnecessary historical reprocessing by ~65% and applied data quality checks (record counts, required fields, and duplicates, consistency) to maintain high data reliability.
- Supported ongoing production monitoring, troubleshooting, and failure handling, contributing to consistent pipeline availability for analytics consumers.

Azure Databricks Data Pipeline

Technologies: Azure Databricks, PySpark, Spark SQL, Delta Lake, Python, SQL

- Developed large-scale data transformation pipelines for multi-TB telemetry and platform datasets using PySpark and Spark SQL on Azure Databricks.
- Implemented partitioning and incremental processing patterns; built and scheduled 10+ Databricks workflows with validation, logging, schema handling, and failure management.
- Supported production pipelines and resolved data-processing issues in a high-volume environment, ensuring reliable delivery to analytics and engineering teams.

EDUCATION

Master of Technology (M.Tech) – Computer Science & Engineering

Jawaharlal Nehru Technological University (JNTU), Hyderabad | 2024 – 2026

CERTIFICATIONS

- AWS Certified Solutions Architect – Associate (SAA-C03)
- Microsoft Certified: Azure Fundamentals (AZ-900)
- Google Cloud Associate Data Practitioner

AWARDS

Wipro Star Performer Award

Recognized for on-time delivery and quality execution of Big Data / data-platform project work.